

AMATH 483 / 583 (Roche) - Homework Set 4

Due Friday May 2, 5pm PT

April 25, 2025

Homework 4 (90 points)

1. (+10) Given matrix $A = \begin{pmatrix} 1 & 2 \\ 0 & 2 \end{pmatrix}$, evaluate the action of A on the unit balls of \mathbb{R}^2 defined by the 1-norm, 2-norm, and ∞ -norm (induced matrix norms). Submit your work and drawings.
2. (+20) Compiler Optimization of Matrix Multiplication Loop Permutations. Implement C++ templated `gemm`, $C \leftarrow \alpha AB + \beta C$ ($A \in \mathbb{T}^{m \times p}$, $B \in \mathbb{T}^{p \times n}$, $C \in \mathbb{T}^{m \times n}$, $\alpha, \beta \in \mathbb{T}$) for $\{kij\}$ and $\{jki\}$ loop permutations using the specifications provided here:

```
• template<typename T>
  void mm_jki(T a, const std::vector<T>& A, const std::vector<T>& B, T b,
    std::vector<T>& C, int m, int p, int n);

• template<typename T>
  void mm_kij(T a, const std::vector<T>& A, const std::vector<T>& B, T b,
    std::vector<T>& C, int m, int p, int n);
```

You will explore matrix multiply performance applying compiler optimization levels `-O0` and `-O3` for square matrices of dimension $n = 2$ to $n = 512$, stride one to these functions. For reference, recall the loop order for $C_{i,j} \leftarrow \sum_{k=0}^{p-1} A_{i,k} * B_{k,j} \forall [i = 0, \dots, m-1][j = 0, \dots, n-1]$ is $\{ijk\}$, i outer loop, j middle loop, k inner loop. Let each n be measured $ntrial$ times and plot the average performance in FLOPs (flop count / time(seconds)) for each case versus n , $ntrial \geq 3$. Submit plots for both permutation variants that include both FP32 (*float*) and FP64 (*double*) compiler optimization results.

3. (+20) **Row major Matrix class.** Reference the file `matrix_class.hpp` for the starter code for a Matrix class template for row major index referencing and `std::vector<T>` for the matrix storage scheme. Please put your function implementations in file `matrix_class.hpp`, and submit `matrix_class.hpp`.

- (+5) Matrix transpose for $A \in \mathbb{R}^{m \times n}$ is defined $A_{i,j}^T = A_{j,i}$ and so $A^T \in \mathbb{R}^{n \times m}$. This method returns a matrix as defined by the class.

```
-      Matrix<T> transpose() const{}
```

- (+5) Matrix infinity norm for $A \in \mathbb{R}^{m \times n}$ is defined $\|A\|_\infty = \max_{0 \leq i \leq m} \sum_{j=1}^{j=n} |A_{i,j}|$. This method returns a number.

```
-      T infinityNorm() const{}
```

- (+5) Write the method to operator overload multiplication `*` for the matrix class I provided. This method returns a matrix as defined by the class.

```
-      Matrix<T> operator*(const Matrix<T> &other) const{}
```

- (+5) Write the method to operator overload addition `+` for the matrix class I provided. This method returns a matrix as defined by the class.

```
-      template <typename T>
      Matrix<T> Matrix<T>::operator+(const Matrix<T> &other) const{}
```

4. (+10) **Extremum.** Consider the surface defined by $xy + 2xz = 5\sqrt{5}$ and $(x, y, z) \in \mathbb{R}^3$. Find (a) the coordinate instance(s) affiliated with the minimum distance from a point on the surface to the origin, and (b) the value of the minimum distance. You will find (may safely assume) the domain $[-3, 3] \times [-3, 3] \times [-3, 3] \in \mathbb{R}^3$ holds the correct coordinate instance(s).
5. (+10) **IO bandwidth.** Write a C++ function that **writes** type **double** square matrices in column major order to file in binary. Measure the time required to complete the **write** for matrices of dimension 32, 64, 128, ... 16384 (2GB). Write a C++ function that **reads** binary matrices from file to type **double** matrices in memory. Measure the time required to complete each **read** for the same dimensions. (a) Make a single plot of the *read* and *write* measurements with the bandwidth (bytes per second) on the y-axis, and the problem dimension on the x-axis. Submit your plot.
6. (+20) **File access time.** (a) Write C++ functions for the given function declarations that perform *row* and *column* swap operations on a type **double** matrix stored in a file in column major index order. Test the swapping capabilities for correctness. Put the functions you write in file **file_swaps.hpp**. (b) Conduct a performance test for square matrix dimensions 16, 32, 64, 128, ... 8192, measuring the time required to conduct file-based *row* and *column* swaps separately. Let each operation be measured *ntrial* times, $ntrial \geq 3$. Make a single plot of the *row* and *column* swap average times on the y-axis ($\log_{10}(time)$) and the problem dimension on the x-axis. Submit your header file **file_swaps.hpp** and plot. You may find the code snippet helpful.

```

void swapRowsInFile(std::fstream &file, int nRows, int nCols, int i, int j);
void swapColsInFile(std::fstream &file, int nRows, int nCols, int i, int j);

// snippet
#include <iostream>
#include <fstream>
#include <vector>
#include <utility>
#include <algorithm>
#include <cstdlib>
#include <ctime>
#include <cstdio>
#include <chrono>
#include "file_swaps.hpp"

int main(int argc, char *argv[])
{
    // Generate the matrix
    std::vector<double> matrix(numRows * numCols);
    // init matrix elements in column major order
    // write the matrix to a file
    std::fstream file(filename, std::ios::out | std::ios::binary);
    file.write(reinterpret_cast<char *>(&matrix[0]), numRows * numCols * sizeof(double));
    file.close();
    // Open the file in read-write mode for swapping
    std::fstream fileToSwap(filename, std::ios::in | std::ios::out | std::ios::binary);
    // Get random indices i and j for row swapping
    // Measure the time required for row swapping using file I/O
    auto startTime = std::chrono::high_resolution_clock::now();
    // Swap rows i and j in the file version of the matrix
    swapRowsInFile(fileToSwap, numRows, numCols, i, j);
    auto endTime = std::chrono::high_resolution_clock::now();
    std::chrono::duration<double> duration = endTime - startTime;
    // Close the file after swapping
    fileToSwap.close();
    //...
    // after each problem size delete the test file
    std::remove(filename.c_str());
    // ...
}

```